

# Cloud Computing Project – Task 1

by University of Illinois at Urbana-Champaign

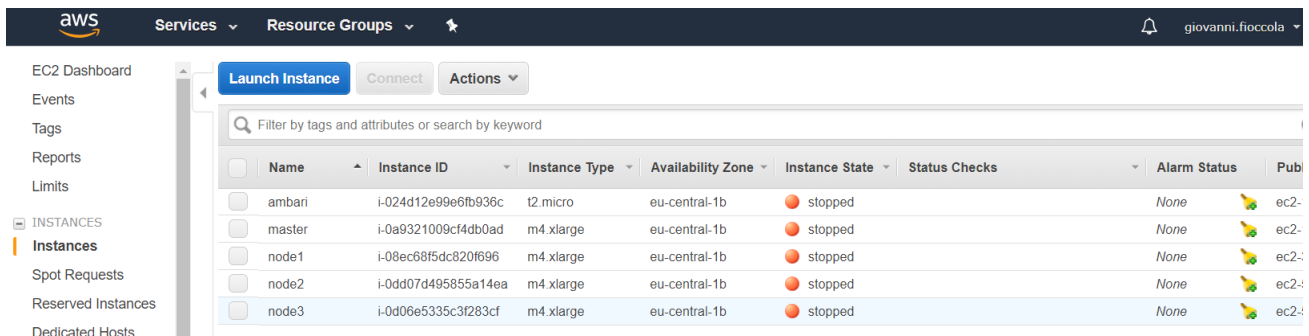
## I. INTRODUCTION

The goal of the Capstone Project is to apply the knowledge acquired through the Cloud computing courses, in order to solve real-world Cloud computing challenges. In particular, the aim of task 1 is to:

- Extract and clean the transportation dataset from the US BTS (Bureau of Transportation Statistics) that is hosted as an Amazon EBS volume snapshot. The dataset contains data and statistics from the US Department of Transportation on aviation, maritime, highway, transit, rail, pipeline, bike/pedestrian and more in CSV format, but only the aviation part of the dataset is useful for this task. Finally, resulting data must be stored in HDFS (Hadoop Distributed File System).
- Answer a set of questions by using Apache Hadoop and/or Spark Streaming, which are two batch processing systems. The results related to questions from Group 2 and question 3.2 are stored in Cassandra.

## II. VIRTUALIZATION ENVIRONMENT

In order to quickly deploy a complete Apache Hadoop stack, Apache Ambari has been used. Ambari is a completely open source management platform for provisioning, managing and monitoring Apache Hadoop clusters by using a graphical user interface. The Apache Hadoop cluster has been deployed on the EC2 Cloud service offered by Amazon Web Services (AWS), and consists of 1 NameNode (Master) and 3 DataNodes (Slaves – Node1, Node2 and Node3). The NameNode is a master server that manages the filesystem namespace and regulates access to files by clients; it executes namespace operations such as opening, closing, and renaming files and directories. The NameNode also determines the mapping of blocks to DataNodes. The DataNodes manage storage attached to the nodes that they run on, and they are responsible for serving read and write requests from the filesystem clients. The DataNodes also perform block creation, deletion, and replication upon instruction from the NameNode. Finally, the Ambari server runs on a fifth EC2 instance.



Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public IP Address
ambari	i-024d12e99e6fb936c	t2.micro	eu-central-1b	stopped	OK	None	ec2-54...
master	i-0a9321009cf4db0ad	m4.xlarge	eu-central-1b	stopped	OK	None	ec2-54...
node1	i-08ec68f5dc820f696	m4.xlarge	eu-central-1b	stopped	OK	None	ec2-54...
node2	i-0dd07d495855a14ea	m4.xlarge	eu-central-1b	stopped	OK	None	ec2-54...
node3	i-0d06e5335c3f283cf	m4.xlarge	eu-central-1b	stopped	OK	None	ec2-54...

Figure 1 - Virtualization environment

## III. DATA ACQUISITION AND FILTERING

A new EBS volume from the existing EBS snapshot (ID snap-e1608d88 for Linux/Unix in the us-east-1 region) of the transportation dataset has been created and attached to the Master EC2 instance. By using the `lsblk` command, it is possible to view the available disk devices and their mount points. A new folder has been created (`/mnt/data`) and the EBS volume has been attached to the filesystem:

```
$ sudo mkdir /mnt/data/
```

```
$ sudo mount /dev/xvdf /mnt/data
```

The attached EBS data volume includes data and statistics from the US Department of Transportation on aviation, maritime, highway, transit, rail, pipeline, bike/pedestrian and more, but only the aviation part of the dataset is useful for this task. Therefore, after mounting and attaching the volume, only the aviation directory has been considered, and more specifically the “airline ontime” data that contains on-time performance for each flight. A python script processes all subdirectories of “airline ontime” and the zip archives included in the subdirectories, ignores the “readme” files and searches for CSV extensions.

```
120     for archived_file in archived_files:
121         logger.debug("Dealing with %s" %(archived_file))
122
123         #determine file type
124         file_type = os.path.splitext(archived_file)[-1]
125
126         #if file is a readme, ignore it
127         if file_type == ".html":
128             logger.debug("ignoring %s" %(archived_file))
129
130         elif file_type == ".csv":
131             logger.debug("extract %s to %s" %(archived_file, workdir))
132             archive.extract(archived_file, workdir)
```

Finally, the script divides input files in chunks and compresses CSV files in “.gz” format by using Pigz, performing a compression through concurrent threads, in order to use multiple processors and cores. Data is sent to HDFS by using put command.

```
143     #Try to split input files in chunks
144     cmd = "split --lines=%s %s %s." %(MAX_LINES, archived_file, basename)
145     cmds = shlex.split(cmd)
146     helper.launch(cmds)
147
148     #now, remove original file
149     os.remove(os.path.join(workdir, data_file))
150
151     #scan directory for csv, load them in HDFS
152     for chunk in os.listdir(workdir):
153         chunk = os.path.join(workdir, chunk)
154
155         #pack csv in gzip format
156         cmd = "pigz --best %s" %(chunk)
157         cmds = shlex.split(cmd)
158         helper.launch(cmds)
159
160         chunk += ".gz"
161
162         #Load data into HDFS:
163         helper.launch(shlex.split("hadoop fs -put %s %s" %(chunk, raw_data_path)))
```

Using mkdir in Apache Hadoop needs the “hadoop file permissions”; hdfs is a user that has permissions to create folders, so the new user is added to hdfs group. Then the folder is created and finally the owner of the folder is changed:

```
$ sudo usermod -aG hdfs $(whoami)
```

```
$ sudo -u hdfs hadoop fs -mkdir -p /user/giovanni/capstone/airline_ontime/raw_data/
```

```
$ sudo -u hdfs hadoop fs -chown -R $(whoami) /user/giovanni/
```

Then the python script is executed:

```
$ python AirlineToHDFS.py --input_path=/mnt/data/aviation/airline_ontime/ \
--output_path=/user/giovanni/capstone/airline_ontime/raw_data/
```

A pig script allows to load data from HDFS, and the loaded data includes the header (the first row). The script gets rid of the header and only the following fields are considered:

FlightDate, AirlineID, FlightNum, Origin, OriginCityName, OriginStateName, Dest, DestCityName, DestStateName, CRSDepTime, DepDelay, CRSArrTime, ArrDelay, Cancelled, CancellationCode, Diverted, CRSElapsedTime, ActualElapsedTime, AirTime, Distance

Finally, filtered data are put into a new folder in HDFS:

```
$ pig -x mapreduce -p input=/user/giovanni/capstone/airline_ontime/raw_data/\
-p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ load_ontime.pig
```

#### IV. SYSTEM INTEGRATION

In order to install Cassandra, the following steps were performed:

- Add the DataStax Community repository to the /etc/apt/sources.list.d/cassandra.sources.list:  
`$ echo "deb http://debian.datastax.com/community stable main" | sudo tee -a /etc/apt/sources.list.d/cassandra.sources.list`
- Add the DataStax repository key to the aptitude trusted keys:  
`$ curl -L http://debian.datastax.com/debian/repo_key | sudo apt-key add -`
- Install the package:  
`$ sudo apt-get update`  
`$ sudo apt-get install dsc20 cassandra=2.0.17`
- Edit /etc/cassandra/cassandra.yaml:  
`$ sudo nano /etc/cassandra/cassandra.yaml`
  - The default cassandra.yaml uses the local, loopback address as its listen and Thrift addresses:  
`listen_address: localhost`  
`rpc_address: localhost`
  - listen\_address and rpc\_address must be set to the IP address of the master:  
`listen_address: 172.31.44.199`  
`rpc_address: 172.31.44.199`
  - Similarly for the the seed information:  
`seeds: "node1,node2,node3"`

The developers of Apache Spark have given thoughtful consideration to Python as a language of choice for data analysis. They have developed the PySpark API for working with RDDs in Python. Furthermore, it supports use of powerful IPythonshell instead of the builtin Python REPL. RDD (Resilient Distributed Dataset) is the basic abstraction in Spark and represents an immutable, partitioned collection of elements that can be operated on in parallel.

The developers of IPython have invested considerable effort in building the IPython Notebook, a system inspired by Mathematica that allows to create “executable documents.” IPython Notebooks can integrate formatted text (Markdown), executable code (Python), mathematical formulae (LaTeX), and graphics/visualizations (matplotlib) into a single document that captures the flow of an exploration and can be exported as a formatted report or an executable script.

The steps followed in order to use iPython with Spark are the following:

- First an IPython profile for use with PySpark has been created:  
`ipython profile create pyspark`

- The file `~/ipython/profile_pyspark/startup/00-pyspark-setup.py` with the following contents has been created:

```
import os
import sys

# Configure the environment
if 'SPARK_HOME' not in os.environ:
    os.environ['SPARK_HOME'] = '/usr/hdp/current/spark-client'

# Create a variable for our root path
SPARK_HOME = os.environ['SPARK_HOME']

# Add the PySpark/py4j to the Python Path
sys.path.insert(0, os.path.join(SPARK_HOME, "python", "lib"))
sys.path.insert(0, os.path.join(SPARK_HOME, "python"))

# Described here: http://blog.cloudera.com/blog/2014/08/how-to-use-ipython-notebook-
with-apache-spark/
sys.path.insert(0, os.path.join(SPARK_HOME, "python/lib/py4j-0.8.2.1-src.zip"))
execfile(os.path.join(SPARK_HOME, 'python/pyspark/shell.py'))
```

- An IPython shell with the profile we just created is started:

```
$ ipython --profile spark
```

Finally, `pyspark-cassandra` is installed. This module provides python support for Apache Spark's Resilient Distributed Datasets from Apache Cassandra CQL rows using Cassandra Spark Connector within PySpark, both in the interactive shell and in python programmes submitted with `spark-submit`.

All MapReduce jobs are submitted to YARN, in order to prepare a dataset that can be processed by Spark. Finally, the Spark Cassandra Connector package allows to move loaded DataFrames from Spark to Apache Cassandra.

## V. QUESTION 1.1

A pig script allows to load data from the filtered transportation dataset. Then, tuples that have the same **Origin** key are grouped together, and tuples that have the same **Dest** key are grouped together too. The number of elements that have the same **Origin** key are counted, the number of elements that have the same **Dest** key are counted too, and results are put in the same table. These results are grouped by airport code and the occurrences are summed up:

```
Total_counts = FOREACH United_group GENERATE FLATTEN($0), SUM(United_counts.count) AS count;
```

Finally, airports are ordered by occurrences thanks to multiple concurrent reduce tasks:

```
$ Popular = ORDER Total_counts BY count DESC PARALLEL 4;
```

This is the pig script:

```
$ pig -x mapreduce -p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ get_top10.pig
```

## VI. QUESTION 1.2

A pig script allows to load data from the filtered transportation dataset. Then, cancelled and diverted flights are excluded. The tuples that have the same **AirlineID** key are grouped together, and for each airline the average arrival delay is computed:

```
airline_delays = FOREACH airlines GENERATE FLATTEN($0), AVG(flights.ArrDelay) AS avg_delay;
```

Finally, airlines are ordered by arrival delays thanks to multiple concurrent reduce tasks.

This is the pig script:

```
$ pig -x mapreduce -p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ top10_airlines.pig
```

## VII. QUESTION 2.1

A pig script allows to load data from the filtered transportation dataset. The airlines lookup dictionary, previously downloaded ([http://www.transtats.bts.gov/Download\\_Lookup.asp?Lookup=L\\_AIRLINE\\_ID](http://www.transtats.bts.gov/Download_Lookup.asp?Lookup=L_AIRLINE_ID)) and put in hadoop filesystem, is loaded. A replicated join is performed:

```
lookup = LOAD '$lookup' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'NOCHANGE', 'SKIP_INPUT_HEADER') AS (AirlineIDsec:chararray,Airline:chararray);
```

```
joined = JOIN filtered BY AirlineID, lookup BY AirlineIDsec USING 'replicated';
```

Then, cancelled and diverted flights are excluded. Only the following fields are considered: **Origin**, **AirlineID**, **Airline**, **DepDelay**:

```
flights = FOREACH arrived_flights GENERATE Origin, AirlineID, Airline, DepDelay;
```

The tuples that have the same **Origin**, **AirlineID** and **Airline** keys are grouped together, and then the average departure delay for each carrier is computed:

```
carrier_delays = FOREACH carrier GENERATE group.Origin, group.AirlineID, group.Airline, AVG(flights.DepDelay) AS avg_delay;
```

The **carrier\_delays** tuples that have the same **Origin** key are grouped together. Finally, for each origin, carriers are ordered by their average departure delays by means of concurrent multiple reduce tasks. Data are then stored into Cassandra.

This is the pig script:

```
$ pig -x mapreduce -p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ -p  
lookup=/user/giovanni/capstone/lookup/Lookup_AirlineID.csv -p  
results=/user/giovanni/capstone/airline_ontime/top10_carriersByAirport/ top10_carriersByAirport.pig
```

Cassandra table definition is the following:

```
CREATE KEYSPACE capstone WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

```
USE capstone;
```

```
CREATE TABLE carriersbyairport ( origin TEXT, airlineid INT, airline TEXT, depdelay FLOAT, PRIMARY  
KEY(origin, depdelay));
```

To show the results, the following script is used:

```
$ cqlsh -f top10_carriersByAirport.cql 172.31.44.199
```

It includes for example:

```
SELECT origin, airlineid, depdelay, airline FROM carriersbyairport WHERE origin = 'CMI';)
```

## VIII. QUESTION 2.2

A pig script allows to load data from the filtered transportation dataset. Then, cancelled and diverted flights are excluded. Only the following fields are considered: **Origin**, **Dest**, **DepDelay**:

```
flights = FOREACH arrived_flights GENERATE Origin, Dest, DepDelay;
```

The tuples that have the same **Origin** and **Dest** keys are grouped together, and then the average departure delay for each path is computed:

```
path_delays = FOREACH path GENERATE group.Origin, group.Dest, AVG(flights.DepDelay) AS avg_delay;
```

The **path\_delays** tuples that have the same **Origin** key are grouped together. Finally, for each origin, **path\_delays** are ordered by their average delays by means of concurrent multiple reduce tasks. Data are then stored into Cassandra.

This is the pig script:

```
$ pig -x mapreduce -p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ \
-p results=/user/giovanni/capstone/airline_ontime/top10_airportsByAirport/ top10_airportsByAirport.pig
```

Cassandra table definition is the following:

```
CREATE KEYSPACE capstone WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

```
USE capstone;
```

```
CREATE TABLE airportsbyairport ( origin TEXT, destination TEXT, depdelay FLOAT, PRIMARY KEY(origin, depdelay));
```

To show the results, the following script is used:

```
$ cqlsh -f top10_airportsByAirport.cql 172.31.44.199
```

It includes for example:

```
SELECT origin, destination, depdelay FROM airportsbyairport WHERE origin = 'CMI';
```

## IX. QUESTION 2.3

A pig script allows to load data from the filtered transportation dataset. The airlines lookup dictionary, previously downloaded ([http://www.transtats.bts.gov/Download\\_Lookup.asp?Lookup=L\\_AIRLINE\\_ID](http://www.transtats.bts.gov/Download_Lookup.asp?Lookup=L_AIRLINE_ID)) and put in hadoop filesystem, is loaded. A replicated join is performed:

```
lookup = LOAD '$lookup' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'NOCHANGE', 'SKIP_INPUT_HEADER') AS (AirlineIDsec:chararray,Airline:chararray);
```

```
joined = JOIN filtered BY AirlineID, lookup BY AirlineIDsec USING 'replicated';
```

Then, cancelled and diverted flights are excluded. Only the following fields are considered: **Origin**, **Dest**, **AirlineID**, **Airline**, **ArrDelay**:

```
flights = FOREACH arrived_flights GENERATE Origin, Dest, AirlineID, Airline, ArrDelay;
```

The tuples that have the same **Origin**, **Dest**, **AirlineID** and **Airline** keys are grouped together, and then the average arrival delay for each carrier is computed:

```
carrier_delays = FOREACH carrier GENERATE group.Origin, group.Dest, group.AirlineID, group.Airline,
AVG(flights.ArrDelay) AS avg_delay;
```

The `carrier_delays` tuples that have the same `Origin` and `Dest` keys are grouped together. Finally, for each origin, carriers are ordered by their average arrival delays by means of concurrent multiple reduce tasks. Data are then stored into Cassandra.

This is the pig script:

```
$ pig -x mapreduce -p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ -p
lookup=/user/giovanni/capstone/lookup/Lookup_AirlineID.csv -p
results=/user/giovanni/capstone/airline_ontime/top10_carriersByPath/top10_carriersByPath.pig
```

Cassandra table definition is the following:

```
CREATE KEYSPACE capstone WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1};
USE capstone;
CREATE TABLE carriersbypath ( origin TEXT, destination TEXT, airlineid INT, airline TEXT, arrdelay FLOAT,
PRIMARY KEY(origin, destination, arrdelay));
```

To show the results, the following script is used:

```
$ cqlsh -f top10_carriersByPath.cql 172.31.44.199
```

It includes for example:

```
SELECT origin, destination, airlineid, arrdelay, airline FROM carriersbypath WHERE origin = 'CMI' AND
destination = 'ORD';
```

## X. QUESTION 3.1

A pig script allows to load data from the filtered transportation dataset. Then, tuples that have the same `Origin` key are grouped together, and tuples that have the same `Dest` key are grouped together too. The number of elements that have the same `Origin` key are counted, the number of elements that have the same `Dest` key are counted too, and results are put in the same table. These results are grouped by airport code and the occurrences are summed up. Finally, airports are ordered by occurrences thanks to multiple concurrent reduce tasks:

```
$ Popular = ORDER Total_counts BY count DESC PARALLEL 4;
```

The RANK operator works with the `count` field. It sorts the relation on this field and prepends the rank value to each tuple. The rank of a tuple is one plus the number of different rank values preceding it. DESC sorts in descending order.

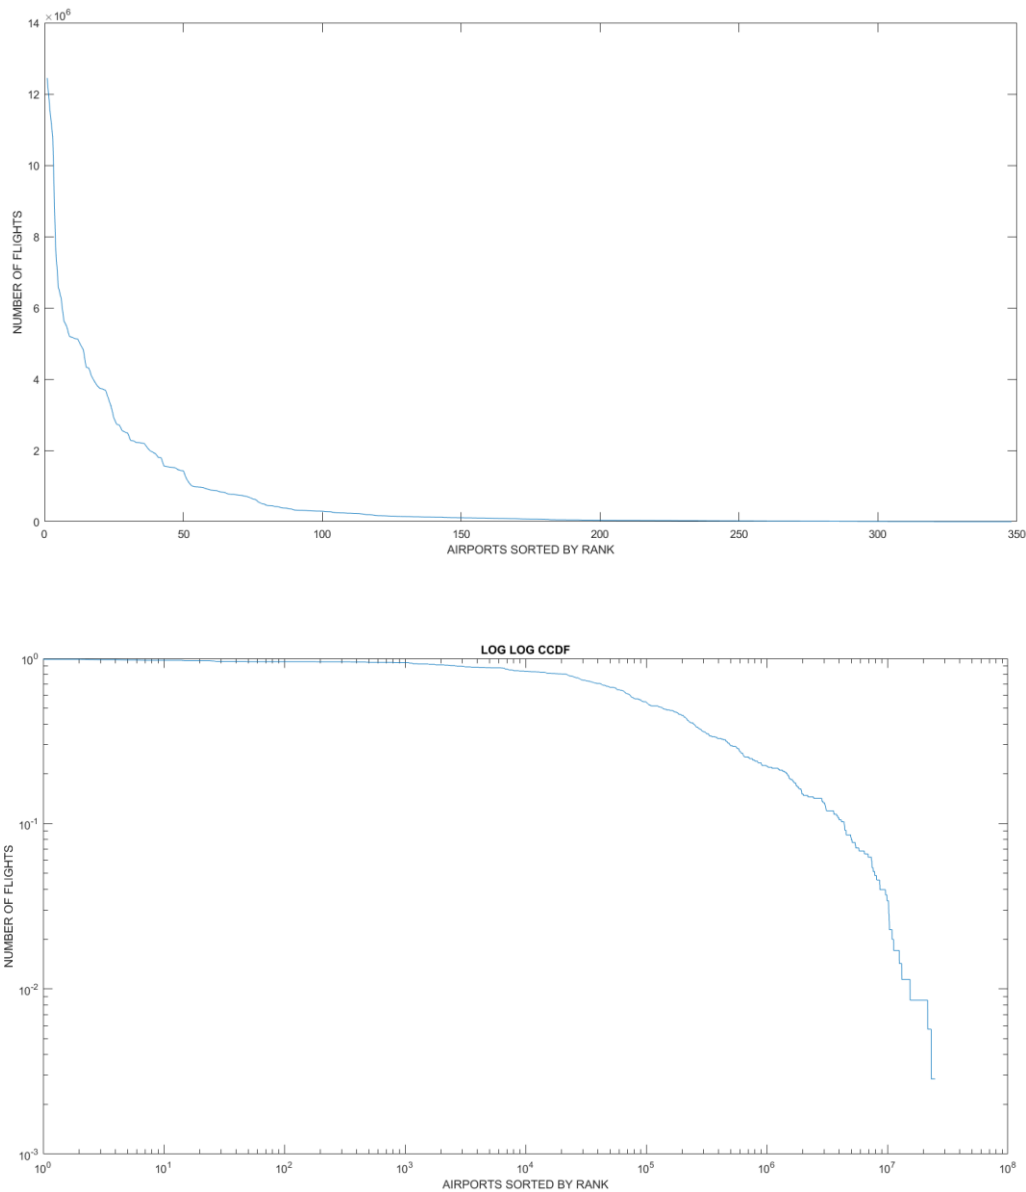
```
Ranked = RANK Popular BY count DESC;
```

This is the pig script:

```
$ pig -x mapreduce -p filtered=/user/giovanni/capstone/airline_ontime/filtered_data/ -p
popular=/user/giovanni/capstone/airline_ontime/popular/ get_popular.pig
```

Zipf's law is an empirical law formulated by using mathematical statistics, and it refers to the fact that many types of data, studied in the physical and social sciences, can be approximated with a Zipfian distribution, one of a family of related discrete power law probability distributions. Zipf's law states that, given some corpus of natural language utterances, the frequency of any word is inversely proportional to its rank in the

frequency table. Therefore, the airport with a higher rank should have a double number of flights with respect to the next airport with lower rank. The CCDF of the popularity for airports should be a straight line if it had been a Zipf distribution. Instead, the popularity of the airports follows a lognormal distribution.



## XI. QUESTION 3.2

A python script allows to read the CSV Data of filtered transportation dataset into an RDD:

```
ontime_data = sc.textFile(DATA_DIR).map(split).map(parse)
```

Then, cancelled and diverted flights are excluded:

```
arrived_data = ontime_data.filter(lambda x: x.Cancelled is False and x.Diverted is False and x.FlightDate.year == 2008 and x.CRSDepTime is not None and x.CRSArrTime is not None and x.ArrDelay is not None)
```

In order to extract date, departure/arrival time scheduled, and the delay, only the following fields are considered:



```
FlightData = arrived_data.map(lambda m: (m.FlightDate, m.Origin, m.Dest, m.FlightNum, m.CRSDepTime,
m.CRSArrTime, m.ArrDelay))
```

Data are filtered because Tom wants his flights scheduled to depart airport X before 12:00 PM local time and to depart airport Y after 12:00 PM local time:

```
path1 = FlightData.filter(lambda (flightdate, origin, dest, flightnum, crsdeptime, crsarrrtime, arrdelay):
crsdeptime < time(hour=12, minute=00))
```

```
path2 = FlightData.filter(lambda (flightdate, origin, dest, flightnum, crsdeptime, crsarrrtime, arrdelay):
crsdeptime > time(hour=12, minute=00))
```

It is possible to transform path by **Origin** and **Dest** keys, in order to join path1 destination with path2 origin. Since the the second leg of the journey (flight Y-Z) must depart two days after the first leg (flight X-Y), it is possible to subtract two days from the second dataset, then join by two keys:

```
destPath1 = path1.keyBy(lambda (flightdate, origin, dest, flightnum, crsdeptime, crsarrrtime, arrdelay): (dest,
flightdate))
```

```
originPath2 = path2.keyBy(lambda (flightdate, origin, dest, flightnum, crsdeptime, crsarrrtime, arrdelay):
(origin, flightdate-timedelta(2)))
```

It is possible to do a join between the two paths. The airport Y (path1 dest, path2 origin) is the key. When a join is performed, RDD with the same keys are located on the same nodes. Calling a map to transform values could be inefficient since the new keys will invalidate partitioning by keys. So, dates are filtered out before calling a new key map phase:

```
joinedPath = destPath1.join(originPath2).values()
```

Tom wants to arrive at each destination with as little delay as possible, so it is possible to sum delays for each of the two paths, then it is possible to order by such values:

```
twoDaysPathFlat = joinedPath.map(lambda ((flightdate1, origin1, dest1, flightnum1, crsdeptime1,
crsarrrtime1, arrdelay1), (flightdate2, origin2, dest2, flightnum2, crsdeptime2, crsarrrtime2, arrdelay2)):
((flightdate1, flightdate2, origin1, dest1, dest2), [(flightdate1, origin1, dest1, flightnum1, crsdeptime1,
crsarrrtime1, arrdelay1, flightdate2, origin2, dest2, flightnum2, crsdeptime2, crsarrrtime2, arrdelay2,
arrdelay1+arrdelay2)]))
```

First, flights from the same dataset with the same path are summed up. aggregatebykey first sums partitions on the same node, then intra-nodes:

```
twoDaysPathSorted = twoDaysPathFlat.aggregateByKey([], getBest, getBest)
```

Finally, data are stored into Cassandra.

The PYSPARK\_CASSANDRA variable has been set to the pyspark directory:

```
PYSPARK_CASSANDRA=/home/giovanni/capstone/pyspark-cassandra/target
```

Then jar files are added:

```
export PYSPARK_SUBMIT_ARGS="--jars ${PYSPARK_CASSANDRA}/pyspark_cassandra-0.1.5.jar \
--driver-class-path ${PYSPARK_CASSANDRA}/pyspark_cassandra-0.1.5.jar \
--py-files ${PYSPARK_CASSANDRA}/pyspark_cassandra-0.1.5-py2.7.egg \
--conf spark.cassandra.connection.host=172.31.44.199"
```

The pyspark\_cassandra library has been imported in order to use Cassandra:

```
import pyspark_cassandra
```

The application is launched by using the spark-submit command:

```
spark-submit $PYSPARK_SUBMIT_ARGS --master yarn --executor-cores=3 --num-executors 3 --driver-memory=3G --executor-memory=4G best2pathsnoakafka.py
```

Cassandra table definition is the following:

```
CREATE KEYSPACE capstone WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

```
USE capstone;
```

```
CREATE TABLE best2path (startdate TEXT, flightnum1 INT, origin1 TEXT, dest1 TEXT, departure1 TIMESTAMP, arrival1 TIMESTAMP, arrdelay1 FLOAT, flightnum2 INT, origin2 TEXT, dest2 TEXT, departure2 TIMESTAMP, arrival2 TIMESTAMP, arrdelay2 FLOAT, PRIMARY KEY(origin1, dest1, dest2, startdate));
```

To show the results, the following script is used:

```
$ cqlsh -f best2paths.cql 172.31.44.199
```

It includes for example:

```
18 /* CMI → ORD → LAX, 04/03/2008 */
19
20 SELECT startdate,
21         flightnum1,
22         origin1,
23         dest1,
24         departure1,
25         arrivall,
26         arrdelay1,
27         flightnum2,
28         origin2,
29         dest2,
30         departure2,
31         arrival2,
32         arrdelay2
33 FROM best2path
34 WHERE origin1 = 'CMI' AND
35        dest1 = 'ORD' AND
36        dest2 = 'LAX' AND
37        startdate = '04/03/2008';
38
```

## XII. OPTIMIZATIONS

The MapReduce algorithm has two important tasks: Map and Reduce. Map takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (key/value pairs). On the other hand, Reduce takes the output from a map as an input and combines the data tuples into smaller set of tuples. In MapReduce, the data is distributed over the cluster and processed. The difference in Spark is that it performs in-memory processing of data. This in-memory processing is a faster process as there is no time spent in moving the data/processes in and out of the disk, whereas MapReduce requires a lot of time to perform these input/output operations, thereby increasing latency. So Spark is used when a large set of files needs to be processed.

In HDFS, the zip files are converted into plain-text CSV files because zip files cannot be split, while CSV files are split into blocks and multiple concurrent map tasks can read the same file that is distributed among different nodes.

# RESULTS

## RESULTS 1.1

```

ubuntu@ip-172-31-44-199: /svn/repos
2017-10-27 19:49:06,050 [main] INFO org.apache.hadoop.yarn.client.AHSPProxy - Connecting to Application History server at ip-172-31-44-199.eu-central-1.compute.internal/172.31.44.199:10200
2017-10-27 19:49:06,051 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server
2017-10-27 19:49:06,073 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2017-10-27 19:49:06,075 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2017-10-27 19:49:06,077 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-10-27 19:49:06,077 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(ORD,12449354)
(ATL,11540422)
(DFW,10799303)
(LAX,7723596)
(PHX,6585534)
(DEN,6273787)
(DTW,5636622)
(IAH,5480734)
(MSP,5199213)
(SFO,5171023)
2017-10-27 19:49:06,132 [main] INFO org.apache.pig.Main - Pig script completed in 11 minutes, 40 seconds and 754 milliseconds (700754 ms)
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime$

```

## RESULTS 1.2

```

ubuntu@ip-172-31-44-199: /svn/repos
2017-10-27 20:02:46,541 [main] INFO org.apache.hadoop.yarn.client.AHSPProxy - Connecting to Application History server at ip-172-31-44-199.eu-central-1.compute.internal/172.31.44.199:10200
2017-10-27 20:02:46,543 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server
2017-10-27 20:02:46,565 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2017-10-27 20:02:46,570 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2017-10-27 20:02:46,573 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-10-27 20:02:46,573 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(19690,-1.01180434574519)
(19678,1.1569234424812056)
(19391,1.4506385127822803)
(20295,4.747609195734892)
(20384,5.3224309999287875)
(20436,5.465881148819851)
(19386,5.557783392671835)
(19393,5.5607742598815735)
(20304,5.736312463662878)
(20363,5.8671846616957595)
2017-10-27 20:02:46,651 [main] INFO org.apache.pig.Main - Pig script completed in 7 minutes, 29 seconds and 994 milliseconds (449994 ms)
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime$

```

## RESULTS 2.1

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime
```

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime$ cqlsh -f top10_carriersByAirport.cql 172.31.44.199
2022 rows imported in 0.847 seconds.
```

origin	airlineid	depdelay	airline
CMI	20417	0.612	Comair Inc.: OH (1)
CMI	20355	1.8244	US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
CMI	20211	3.7446	Trans World Airways LLC: TW
CMI	19822	4.3699	Piedmont Aviation Inc.: PI
CMI	20404	6.0802	Independence Air: DH
CMI	20366	6.6651	ExpressJet Airlines Inc.: EV
CMI	20398	7.9928	Envoy Air: MQ

(7 rows)

origin	airlineid	depdelay	airline
BWI	20436	0.7469	Frontier Airlines Inc.: F9
BWI	20384	4.7619	Pan American World Airways (1): PA (1)
BWI	19704	5.1138	Continental Air Lines Inc.: CO
BWI	20378	5.5047	Mesa Airlines Inc.: YV
BWI	19386	5.6557	Northwest Airlines Inc.: NW
BWI	19805	5.9134	American Airlines Inc.: AA
BWI	20363	7.2207	Endeavor Air Inc.: 9E
BWI	20355	7.4617	US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
BWI	19790	7.6411	Delta Air Lines Inc.: DL
BWI	19977	7.6708	United Air Lines Inc.: UA

(10 rows)

origin	airlineid	depdelay	airline
MIA	20363	-3.6667	Endeavor Air Inc.: 9E
MIA	20366	1.2026	ExpressJet Airlines Inc.: EV
MIA	20312	1.7822	ATA Airlines d/b/a ATA: TZ
MIA	20374	1.8627	ExpressJet Airlines Inc. (1): XE
MIA	20384	4.0265	Pan American World Airways (1): PA (1)
MIA	19386	4.4239	Northwest Airlines Inc.: NW
MIA	20355	6.0187	US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
MIA	19977	6.8083	United Air Lines Inc.: UA
MIA	20295	7.5121	Midway Airlines Inc. (1): ML (1)
MIA	19822	7.9504	Piedmont Aviation Inc.: PI

(10 rows)

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime
```

origin	airlineid	depdelay	airline
LAX	20398	2.4	Envoy Air: MQ
LAX	20304	4.2141	SkyWest Airlines Inc.: OO
LAX	20437	4.6878	AirTran Airways Corporation: FL
LAX	20312	4.7589	ATA Airlines d/b/a ATA: TZ
LAX	19391	4.8054	Pacific Southwest Airlines: PS (1)
LAX	19386	5.0929	Northwest Airlines Inc.: NW
LAX	20436	5.7212	Frontier Airlines Inc.: F9
LAX	19690	5.8176	Hawaiian Airlines Inc.: HA
LAX	20378	6.0363	Mesa Airlines Inc.: YV
LAX	20355	6.7294	US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)

(10 rows)

origin	airlineid	depdelay	airline
IAH	19386	3.5156	Northwest Airlines Inc.: NW
IAH	20384	3.9016	Pan American World Airways (1): PA (1)
IAH	19822	3.9522	Piedmont Aviation Inc.: PI
IAH	20355	5.0298	US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
IAH	20436	5.5495	Frontier Airlines Inc.: F9
IAH	19805	5.6169	American Airlines Inc.: AA
IAH	20211	6.0275	Trans World Airways LLC: TW
IAH	19393	6.2131	Southwest Airlines Co.: WN
IAH	20304	6.5833	SkyWest Airlines Inc.: OO
IAH	20398	6.7098	Envoy Air: MQ

(10 rows)

origin	airlineid	depdelay	airline
SFO	20312	3.9351	ATA Airlines d/b/a ATA: TZ
SFO	20398	4.7827	Envoy Air: MQ
SFO	20436	5.1559	Frontier Airlines Inc.: F9
SFO	20384	5.2839	Pan American World Airways (1): PA (1)
SFO	19386	5.7175	Northwest Airlines Inc.: NW
SFO	19391	6.2567	Pacific Southwest Airlines: PS (1)
SFO	19790	6.5125	Delta Air Lines Inc.: DL
SFO	19704	7.0474	Continental Air Lines Inc.: CO
SFO	20355	7.4795	US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
SFO	20211	7.7631	Trans World Airways LLC: TW

(10 rows)

RESULTS 2.2

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime
```

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime$ cqlsh -f top10_airportsByAirport.cql 172.31.44.199
2325 rows imported in 0.687 seconds.
```

```
origin | destination | depdelay
-----+-----+-----
CMI | PIT | 1.1078
CMI | CVG | 1.9021
CMI | DAY | 2.928
CMI | PIA | 3.4126
CMI | STL | 3.8417
CMI | DFW | 5.9781
CMI | ATL | 6.6651
CMI | ORD | 8.1635
```

```
(8 rows)
```

```
origin | destination | depdelay
-----+-----+-----
BWI | MLB | 1.1615
BWI | DAB | 1.5085
BWI | SRQ | 1.55
BWI | IAD | 1.7972
BWI | UCA | 3.3259
BWI | CHO | 3.7449
BWI | BGM | 3.7951
BWI | DCA | 3.9706
BWI | GSP | 4.1928
BWI | SJU | 4.2034
```

```
(10 rows)
```

```
origin | destination | depdelay
-----+-----+-----
MIA | SAN | 1.7104
MIA | BUF | 2
MIA | SLC | 2.4108
MIA | HOU | 2.8885
MIA | ISP | 3.6491
MIA | MEM | 3.6863
MIA | GNV | 3.9607
MIA | PSE | 3.9758
MIA | TLH | 4.2346
MIA | MCI | 4.6122
```

```
(10 rows)
```

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-caps
```

```
origin | destination | depdelay
-----+-----+-----
LAX | SDF | -16
LAX | LAX | -2
LAX | BZN | -0.72727
LAX | MAF | 0
LAX | MFE | 1.1775
LAX | IYK | 1.2698
LAX | MEM | 1.868
LAX | PIE | 1.9262
LAX | SNA | 2.0945
LAX | SBA | 2.2744
```

```
(10 rows)
```

```
origin | destination | depdelay
-----+-----+-----
IAH | MSN | -2
IAH | MLI | -1
IAH | AGS | -0.6173
IAH | EFD | 1.8877
IAH | HOU | 2.172
IAH | JAC | 2.4226
IAH | MTJ | 2.9025
IAH | RNO | 3.2273
IAH | GUC | 3.5374
IAH | BPT | 3.5995
```

```
(10 rows)
```

```
origin | destination | depdelay
-----+-----+-----
SFO | SDF | -10
SFO | MSO | -4
SFO | LGA | -1.7576
SFO | PIE | -1.341
SFO | OAK | -0.81371
SFO | FAR | 0
SFO | BNA | 2.3792
SFO | MEM | 3.2624
SFO | SJC | 3.9856
SFO | MKE | 3.988
```

```
(10 rows)
```

## RESULTS 2.3

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime
```

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime$ cqlsh -f top10_carriersByPath.cql 172.31.44.199
19384 rows imported in 5.607 seconds.
```

```
origin | destination | airlineid | arrdelay | airline
-----+-----+-----+-----+-----
CMI | ORD | 20398 | 10.144 | Envoy Air: MQ
```

```
(1 rows)
```

```
origin | destination | airlineid | arrdelay | airline
-----+-----+-----+-----+-----
IND | CMH | 19704 | -2.5459 | Continental Air Lines Inc.: CO
IND | CMH | 19805 | 5.5 | American Airlines Inc.: AA
IND | CMH | 19991 | 5.6973 | America West Airlines Inc.: HP (Merged with US Airways 9/05.Stopped reporting 10/07.)
IND | CMH | 19386 | 5.7615 | Northwest Airlines Inc.: NW
IND | CMH | 20355 | 6.8785 | US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
IND | CMH | 19790 | 10.688 | Delta Air Lines Inc.: DL
IND | CMH | 19707 | 10.813 | Eastern Air Lines Inc.: EA
```

```
(7 rows)
```

```
origin | destination | airlineid | arrdelay | airline
-----+-----+-----+-----+-----
DFW | IAH | 20384 | -1.5965 | Pan American World Airways (1): PA (1)
DFW | IAH | 20366 | 5.0925 | ExpressJet Airlines Inc.: EV
DFW | IAH | 19977 | 5.4142 | United Air Lines Inc.: UA
DFW | IAH | 19704 | 6.4937 | Continental Air Lines Inc.: CO
DFW | IAH | 20304 | 7.564 | SkyWest Airlines Inc.: OO
DFW | IAH | 20374 | 8.0943 | ExpressJet Airlines Inc. (1): XE
DFW | IAH | 19805 | 8.3812 | American Airlines Inc.: AA
DFW | IAH | 19790 | 8.5985 | Delta Air Lines Inc.: DL
DFW | IAH | 20398 | 9.1032 | Envoy Air: MQ
```

```
(9 rows)
```

```
origin | destination | airlineid | arrdelay | airline
-----+-----+-----+-----+-----
LAX | SFO | 20312 | -7.619 | ATA Airlines d/b/a ATA: TZ
LAX | SFO | 19391 | -2.1463 | Pacific Southwest Airlines: PS (1)
LAX | SFO | 20436 | -2.0287 | Frontier Airlines Inc.: F9
LAX | SFO | 20366 | 6.9646 | ExpressJet Airlines Inc.: EV
LAX | SFO | 19805 | 7.3868 | American Airlines Inc.: AA
LAX | SFO | 20398 | 7.8078 | Envoy Air: MQ
LAX | SFO | 20355 | 7.9647 | US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
LAX | SFO | 19393 | 8.7921 | Southwest Airlines Co.: WN
LAX | SFO | 19704 | 9.3548 | Continental Air Lines Inc.: CO
LAX | SFO | 19386 | 9.8488 | Northwest Airlines Inc.: NW
```

```
(10 rows)
```

```
origin | destination | airlineid | arrdelay | airline
-----+-----+-----+-----+-----
JFK | LAX | 19977 | 3.3139 | United Air Lines Inc.: UA
JFK | LAX | 19991 | 6.6806 | America West Airlines Inc.: HP (Merged with US Airways 9/05.Stopped reporting 10/07.)
JFK | LAX | 19805 | 6.9037 | American Airlines Inc.: AA
JFK | LAX | 19790 | 7.9345 | Delta Air Lines Inc.: DL
JFK | LAX | 20384 | 11.019 | Pan American World Airways (1): PA (1)
JFK | LAX | 20211 | 11.702 | Trans World Airways LLC: TW
```

```
(6 rows)
```

```
origin | destination | airlineid | arrdelay | airline
-----+-----+-----+-----+-----
ATL | PHX | 20437 | 4.5526 | AirTran Airways Corporation: FL
ATL | PHX | 20355 | 6.2881 | US Airways Inc.: US (Merged with America West 9/05. Reporting for both starting 10/07.)
ATL | PHX | 19991 | 8.4814 | America West Airlines Inc.: HP (Merged with US Airways 9/05.Stopped reporting 10/07.)
ATL | PHX | 19707 | 8.9536 | Eastern Air Lines Inc.: EA
ATL | PHX | 19790 | 9.8083 | Delta Air Lines Inc.: DL
```

```
(5 rows)
```

```
ubuntu@ip-172-31-44-199: /home/giovanni/ccc-capstone/ontime$
```

RESULTS 3.1RESULTS 3.2

```

ubuntu@ip-172-31-44-199:/home/giovanni/ccp-capstone/ontime
ubuntu@ip-172-31-44-199:/home/giovanni/ccp-capstone/ontime$
ubuntu@ip-172-31-44-199:/home/giovanni/ccp-capstone/ontime$ cqlsh -f best2paths.cql 172.31.44.199

startdate | flightnum1 | origin1 | dest1 | departure1 | arrival1 | ardelay1 | flightnum2 | origin2 | dest2 | departure2 | arrival2
| ardelay2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
04/03/2008 | 4278 | CMI | ORD | 2008-03-04 07:10:00+0000 | 2008-03-04 08:15:00+0000 | -14 | 607 | ORD | LAX | 2008-03-06 19:50:00+0000 | 2008-03-06 22:20:00+
0000 | -24
(1 rows)

startdate | flightnum1 | origin1 | dest1 | departure1 | arrival1 | ardelay1 | flightnum2 | origin2 | dest2 | departure2 | arrival2
| ardelay2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
09/09/2008 | 845 | JAX | DFW | 2008-09-09 07:25:00+0000 | 2008-09-09 08:55:00+0000 | 1 | 3627 | DFW | CRP | 2008-09-11 16:45:00+0000 | 2008-09-11 18:05:00+
0000 | -7
(1 rows)

startdate | flightnum1 | origin1 | dest1 | departure1 | arrival1 | ardelay1 | flightnum2 | origin2 | dest2 | departure2 | arrival2
| ardelay2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
01/04/2008 | 3755 | SLC | BFL | 2008-04-01 11:00:00+0000 | 2008-04-01 11:53:00+0000 | 12 | 5429 | BFL | LAX | 2008-04-03 14:55:00+0000 | 2008-04-03 15:45:00+
0000 | 6
(1 rows)

startdate | flightnum1 | origin1 | dest1 | departure1 | arrival1 | ardelay1 | flightnum2 | origin2 | dest2 | departure2 | arrival2
| ardelay2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
12/07/2008 | 3534 | LAX | SFO | 2008-07-12 06:50:00+0000 | 2008-07-12 08:05:00+0000 | -13 | 412 | SFO | PHX | 2008-07-14 19:25:00+0000 | 2008-07-14 21:26:00+
0000 | -19
(1 rows)

startdate | flightnum1 | origin1 | dest1 | departure1 | arrival1 | ardelay1 | flightnum2 | origin2 | dest2 | departure2 | arrival2
| ardelay2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
10/06/2008 | 1104 | DFW | ORD | 2008-06-10 07:00:00+0000 | 2008-06-10 09:23:00+0000 | -21 | 2341 | ORD | DFW | 2008-06-12 16:45:00+0000 | 2008-06-12 19:15:00+
0000 | -10

startdate | flightnum1 | origin1 | dest1 | departure1 | arrival1 | ardelay1 | flightnum2 | origin2 | dest2 | departure2 | arrival2
| ardelay2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
01/01/2008 | 944 | LAX | ORD | 2008-01-01 07:05:00+0000 | 2008-01-01 13:02:00+0000 | 1 | 918 | ORD | JFK | 2008-01-03 19:00:00+0000 | 2008-01-03 22:19:00+
0000 | -7
(1 rows)

```